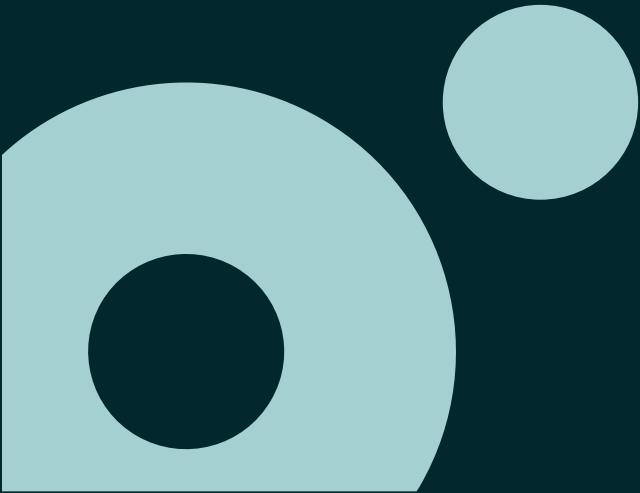
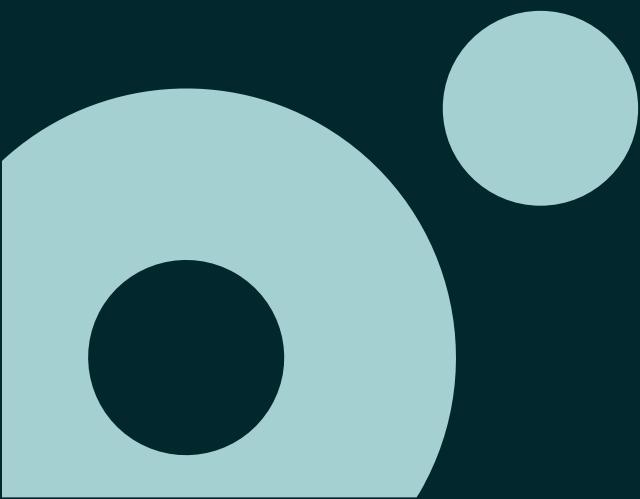


Containers & Kubernetes

Session #07



Services



Motivation

Services

- Kubernetes Pods are created and destroyed to match the state of your cluster making them ephemeral
- Each Pod gets its own IP address, however in a Deployment, the set of Pods running in one moment in time could be different from the set of Pods running that application a moment later
- If some set of Pods provides functionality to other Pods inside your cluster, how do they find out and keep track of which IP address to connect to?

What is a Service?

Services

- An abstraction that defines a logical set of loosely-coupled pods and a policy by which to access them as a network service.
 - Use selectors to define which pods to include.
 - Load balances traffic to Pods (Layer 4)
 - Exposes Pods to other Pods within the cluster
 - Maps external ports to target (container) ports
- Kubernetes automatically updates which Pods are available to which service by creating Endpoints objects.
- 3 types
 - ClusterIP
 - NodePort
 - LoadBalancer

Service Manifest




Service

Service properties

port sets service port
targetPort refers pods port

selector defines pods selector

type defines service type



```
apiVersion: v1
kind: Service
metadata:
  name: sample-svc
spec:
  ports:
    - port: 8080
      targetPort: 80
      name: web
  selector:
    app: sample
    tech: dotnet
  type: ClusterIP
```

Endpoint Object

Service

Endpoint properties

Each block defines pod settings

Each block defines pod settings

Each block defines pod settings

ports defines how to connect to pods

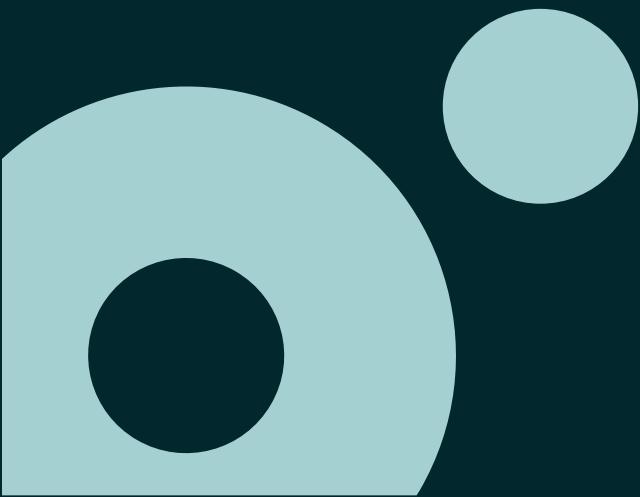
```
apiVersion: v1
kind: Endpoints
metadata:
  name: sample-svc
subsets:
- addresses:
  - ip: 10.1.0.177
    nodeName: docker-desktop
    targetRef:
      kind: Pod
      name: sample-dep-8966bc4c5-67ngv
- ip: 10.1.0.178
    nodeName: docker-desktop
    targetRef:
      kind: Pod
      name: sample-dep-8966bc4c5-nrfpp
- ip: 10.1.0.179
    nodeName: docker-desktop
    targetRef:
      kind: Pod
      name: sample-dep-8966bc4c5-92ts6
ports:
- name: web
  port: 80
  protocol: TCP
```

Load Balancing

Services

- Depends on kube-proxy configuration mode
- User space proxy mode
 - Uses round-robin algorithm to select pods
- Iptables proxy mode (default)
 - Uses random selection of pods
- IPVS proxy mode
 - rr: round-robin
 - lc: least connection (smallest number of open connections)
 - dh: destination hashing
 - sh: source hashing
 - sed: shortest expected delay
 - nq: never queue


ClusterIP



ClusterIP

Services

- Exposes the Service on a cluster-internal IP
- Choosing this value makes the Service only reachable from within the cluster
- Default service type
- Uses service name as a DNS
 - <http://svc-name>
 - <http://svc-name.ns.svc.cluster.local>



Deployment

```
● ● ●

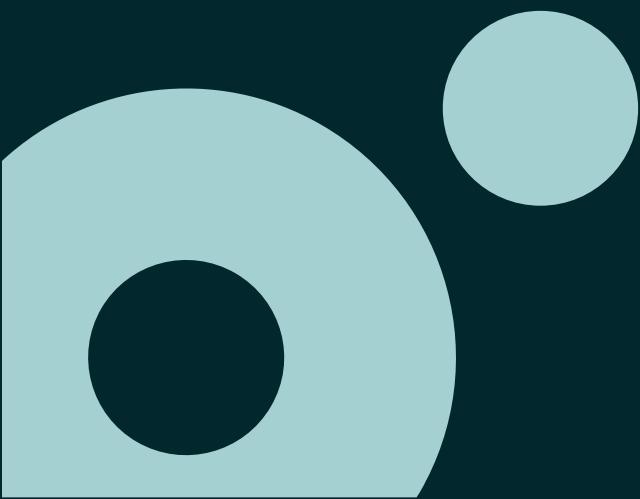
apiVersion: apps/v1
kind: Deployment
metadata:
  name: sample-dep
spec:
  replicas: 3
  selector:
    matchLabels:
      app: sample
      tech: dotnet
  template:
    metadata:
      labels:
        app: sample
        tech: dotnet
    spec:
      containers:
      - name: sample
        image: mcr.microsoft.com/dotnet/samples:aspnetapp
      ports:
      - containerPort: 80
```

ClusterIP



```
apiVersion: v1
kind: Service
metadata:
  name: sample-svc
spec:
  ports:
    - port: 8080
      targetPort: 80
      name: web
  selector:
    app: sample
    tech: dotnet
  type: ClusterIP
```



Node Port



NodePort

Services

- Exposes the Service on each Node's IP at a static port (the NodePort)
- A ClusterIP Service, to which the NodePort Service routes, is automatically created.
- You'll be able to contact the NodePort Service, from outside the cluster, by requesting `<NodeIP>:<NodePort>`
- If **nodePort** property is not specified, is automatically set an port from the range 30000-32767




NodePort



```
apiVersion: v1
kind: Service
metadata:
  name: sample-svc
spec:
  ports:
    - port: 8080
      targetPort: 80
      nodePort: 30000
      name: web
  selector:
    app: sample
    tech: dotnet
  type: NodePort
```



Load Balancer



LoadBalancer

Services

- Exposes the Service externally using a cloud provider's load balancer
- NodePort and ClusterIP Services, to which the external load balancer routes, are automatically created
- On-prem needs manual configuration
- Can be internal (no public IP)




LoadBalancer



```
apiVersion: v1
kind: Service
metadata:
  name: sample-svc
spec:
  ports:
    - port: 8080
      targetPort: 80
      nodePort: 30000
      name: web
  selector:
    app: sample
    tech: dotnet
  type: LoadBalancer
```


Ingress & Ingress Controller



Motivation


Ingress & Ingress Controller

- Creating several Load Balancer services will need several Public Ips to allow access to workloads
- Several issues on this approach
 - Several IPs to handle
 - Security with direct access to the cluster
 - SSL termination
- Needs to have a single ingress (inbound) connection from outside the cluster to inside services

Ingress

Ingress & Ingress Controller


- Exposes HTTP and HTTPS routes from outside the cluster to services within the cluster
- Traffic routing is controlled by rules defined on the Ingress resource



Ingress Controller

Ingress & Ingress Controller

- Kubernetes don't have concrete implementation for this routing
- Ingress are used by Ingress Controllers to implement the routing
- Most used
 - The [NGINX Ingress Controller for Kubernetes](#) works with the [NGINX](#) webserver (as a proxy).
 - The [Traefik Kubernetes Ingress provider](#) is an ingress controller for the [Traefik](#) proxy.



Ingress



```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: simple-fanout-example
spec:
  rules:
  - host: foo.bar.com
    http:
      paths:
      - path: /foo
        pathType: Prefix
        backend:
          service:
            name: service1
          port:
            number: 4200
      - path: /bar
        pathType: Prefix
        backend:
          service:
            name: service2
          port:
            number: 8080
```

